

Istio - Service mesh

Georgios Andrianakis

Senior Software Engineer - Red Hat



We of course need to start with...

Microservices

Microservices are Distributed Systems

They introduce non-trivial complexity into the space between our microservices - the *network*



Microservices Architecture Challenges

- Resilience and fault tolerance
 - Retries
 - Timeouts
 - Circuit Breaker
- Discover, load balance services
- Logging
- Metrics
- Distribute traces

...challenges continued

- Rate Limiting
- Canary Deployments
- A/B Testing
- Fault injection

**How do Microservices Architectures
deal with these challenges?**

The *library* approach

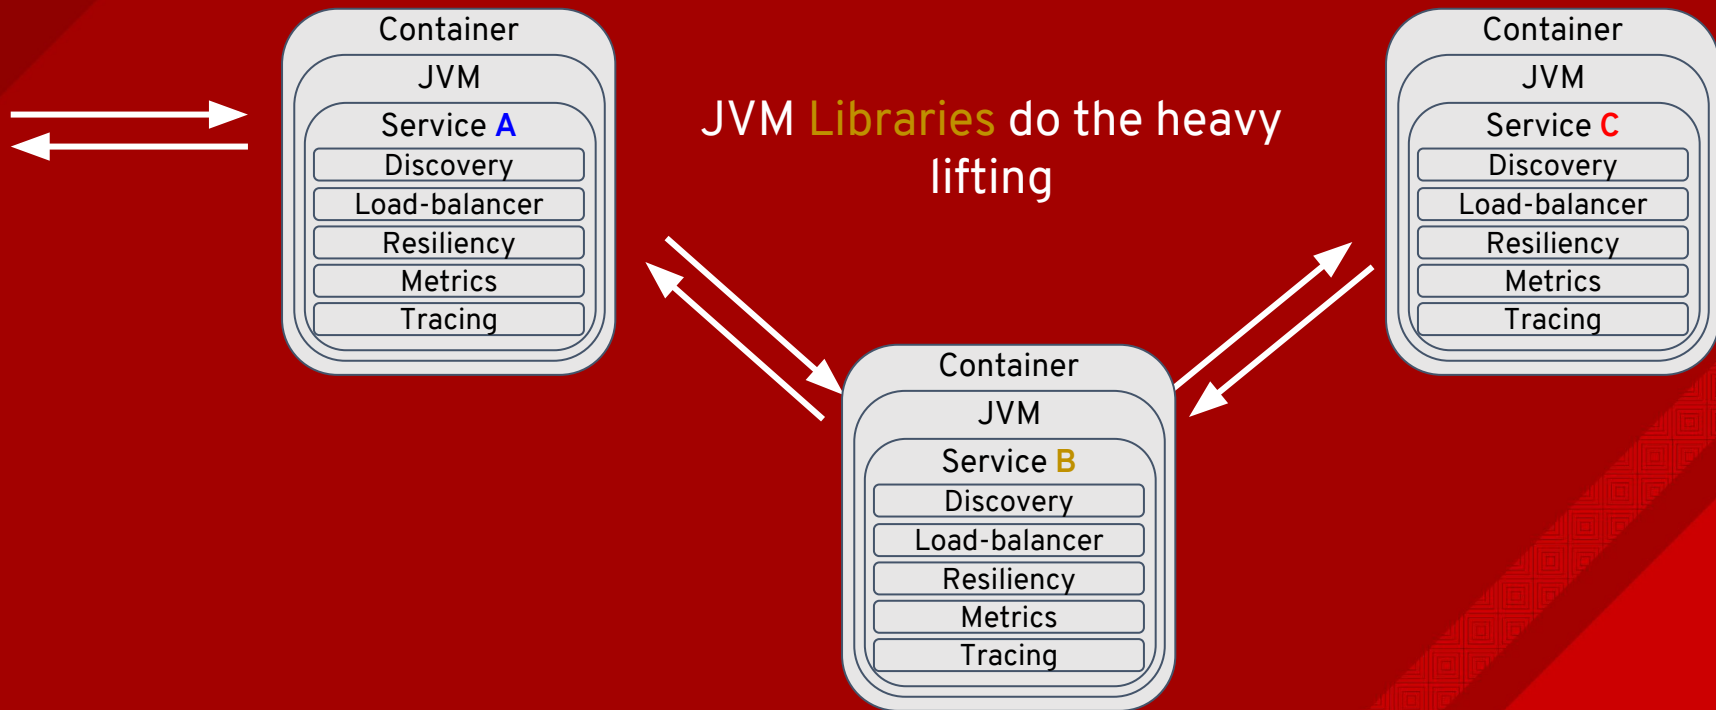


- Netflix Hystrix (circuit breaking / bulk heading)
- Netflix Zuul (edge router)
- Netflix Ribbon (client-side service discovery / load balance)
- Netflix Eureka (service discovery registry)
- Brave / Zipkin (tracing)
- Netflix spectator / atlas (metrics)

But I'm using Spring!!!

- spring-cloud-netflix-hystrix
- spring-cloud-netflix-zuul
- spring-cloud-netflix-eureka-client
- spring-cloud-netflix-ribbon
- spring-cloud-netflix-atlas
-
-
- @Enable....150 different Things

Microservices embedding Capabilities



Drawbacks to the library approach

- Need an implementation for each combination of runtime/framework
 - Might end up forcing specific languages / frameworks on teams
- Need to maintain, upgrade, retire
- Classpath / namespace pollution
 - We end up with large deployables even for very small microservices

The platform approach

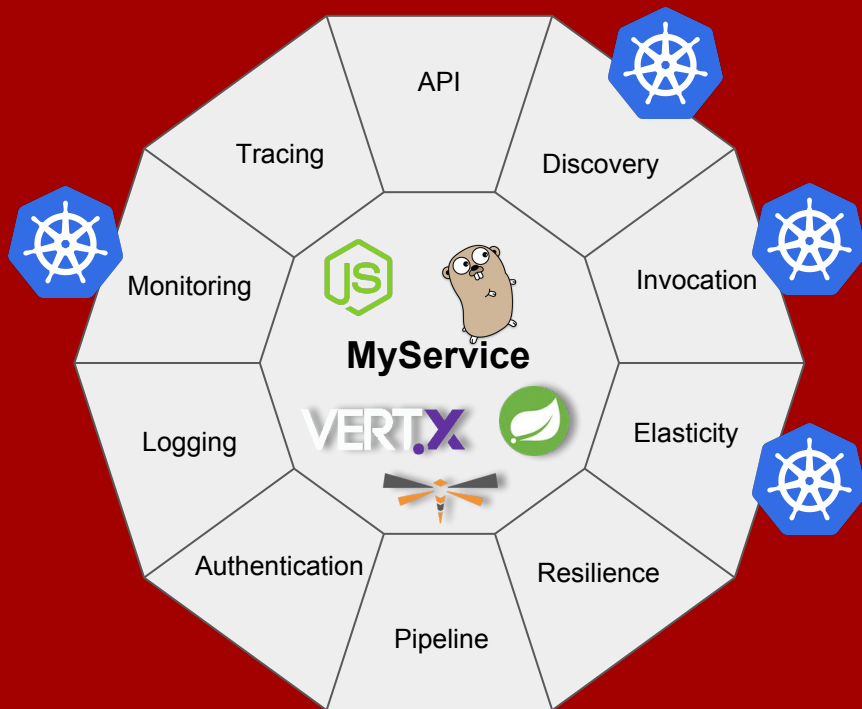
Move horizontal concerns into the platform and apply to all services regardless of implementation.

Benefits of the platform approach

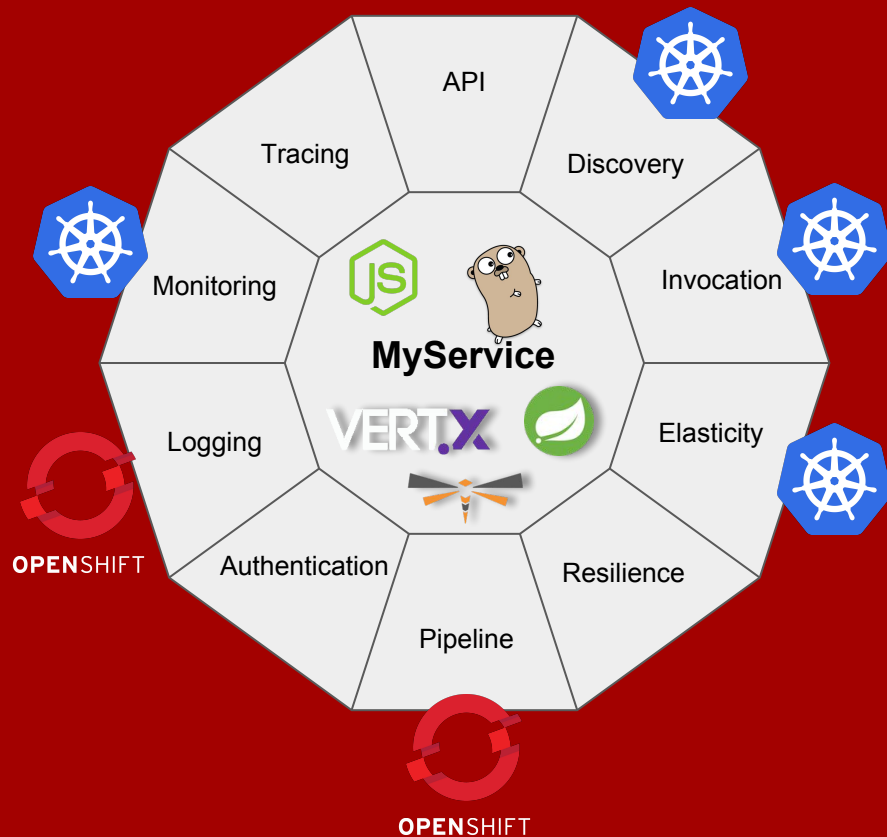
- Allow heterogeneous architectures
- Consistently and correctly enforce approaches to microservices challenges

What do the various platforms look like?

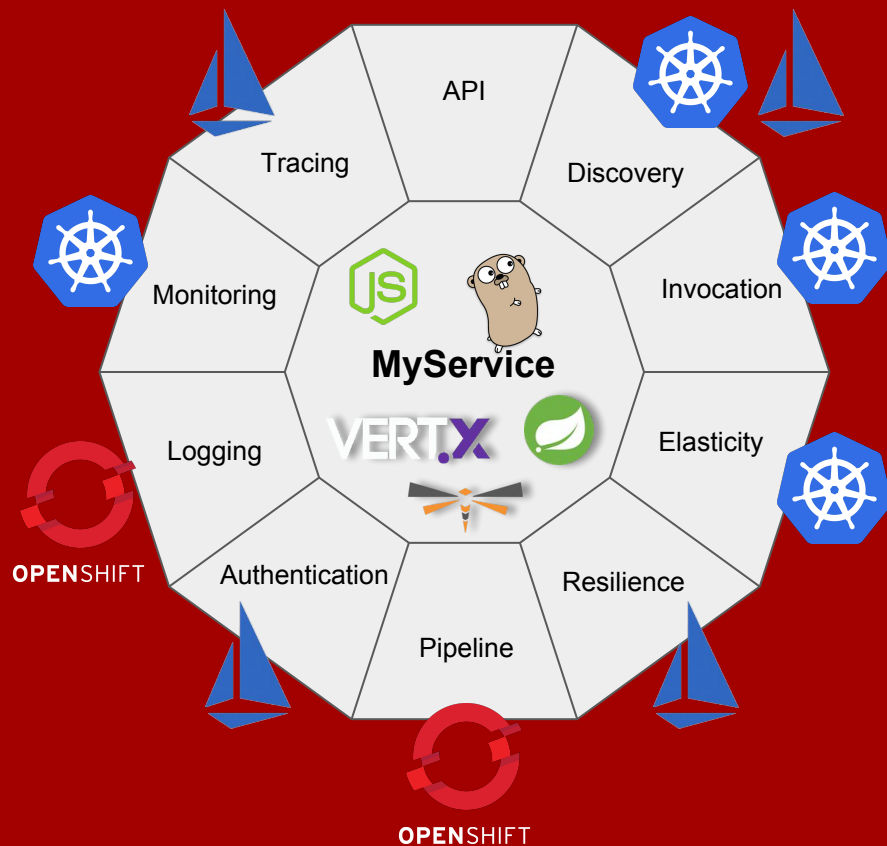
Microservices + Kubernetes



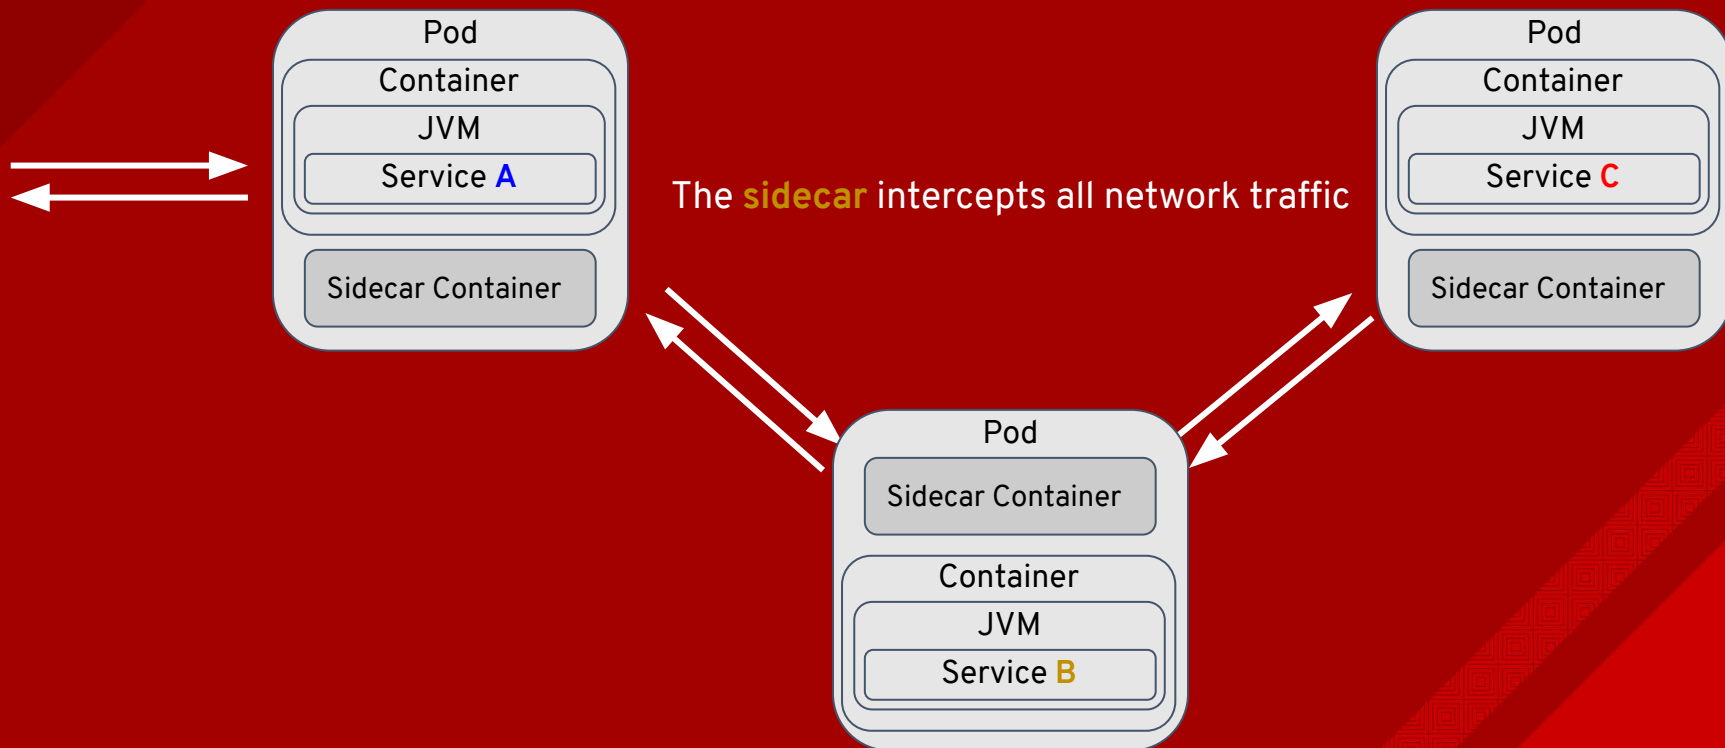
Microservices + OpenShift



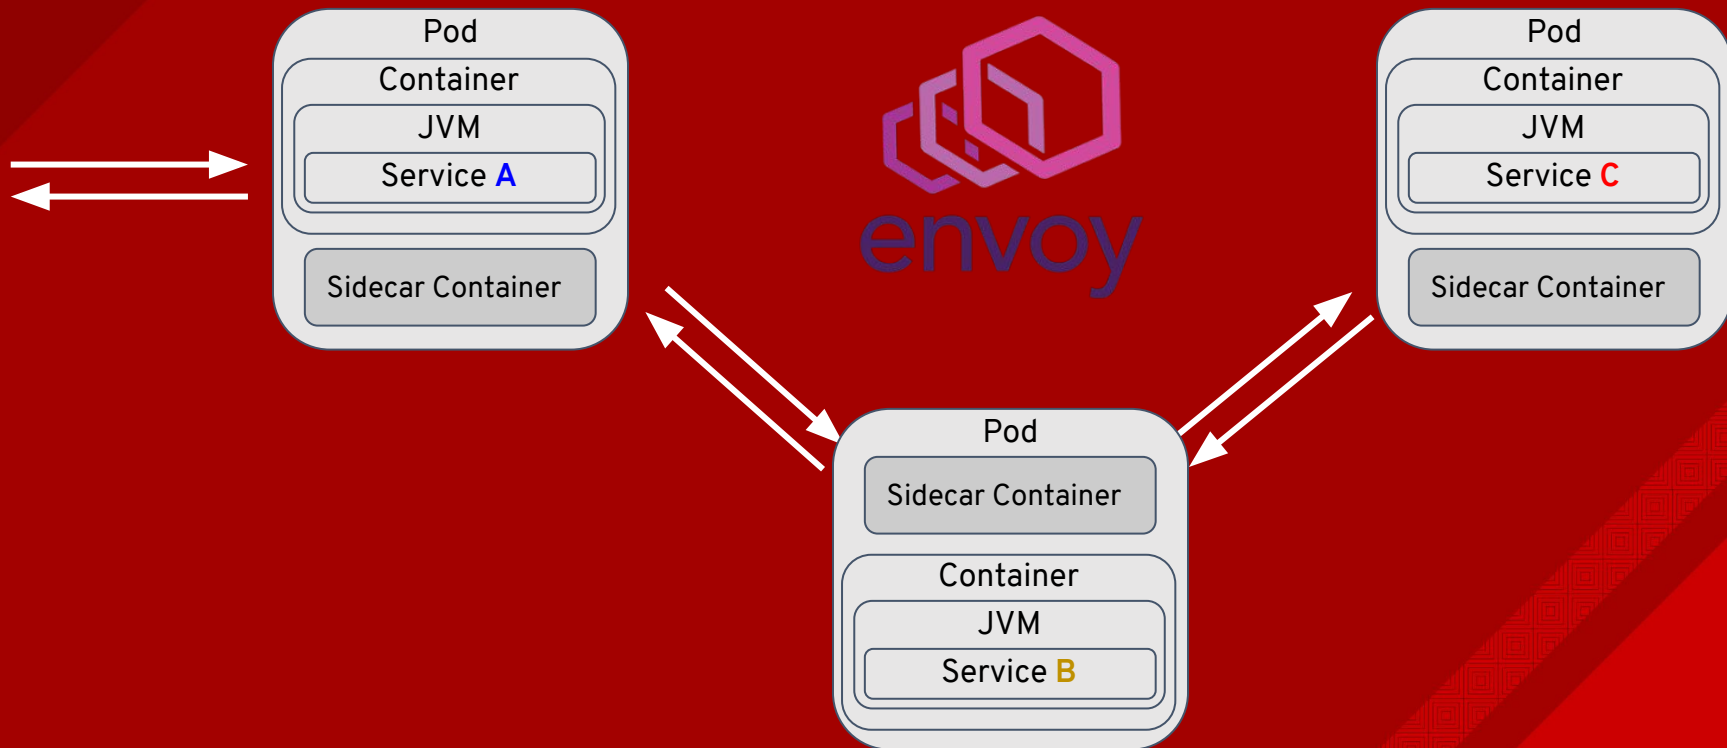
Microservices + Istio



Microservices embedding Capabilities



Envoy is the sidecar in Istio



Meet Envoy Proxy

<http://envoyproxy.io>



ENVOY IS AN OPEN SOURCE EDGE AND SERVICE
PROXY, DESIGNED FOR CLOUD-NATIVE APPLICATIONS

Envoy ...

- Is written in C++, highly parallel, non-blocking
- L3/4 network filter
- out of the box L7 filters
- HTTP 2, including gRPC
- baked in service discovery / health checking
- advanced load balancing
- stats, metrics, tracing

How do we reason about a fleet of
these service proxies in a large cluster?

Time for definitions

A **service mesh** is *decentralized* application-networking infrastructure *between your services* that provides *resilience, security, observability,* and routing *control*.

A service mesh is comprised of a *data plane* and a *control plane*.


Meet Istio.io

<http://istio.io>

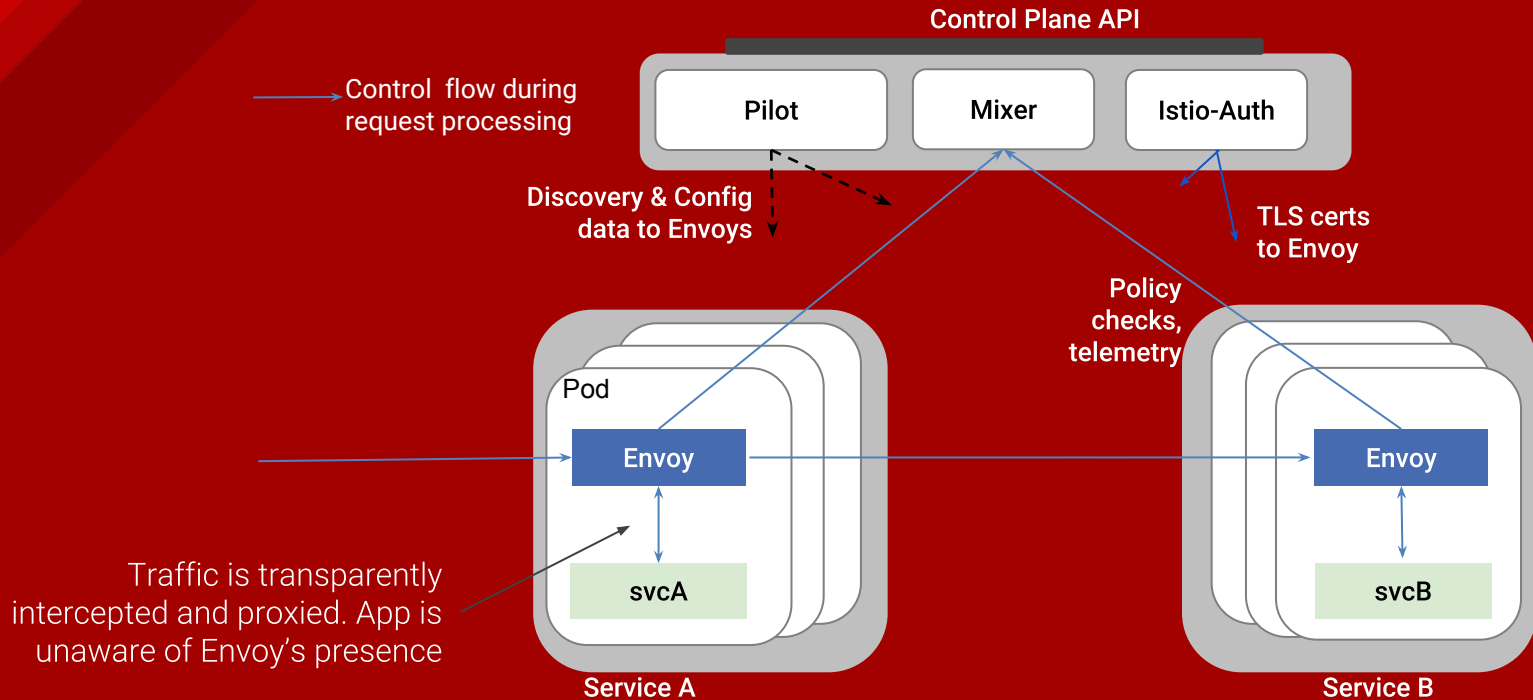


A control plane for service proxies

Components




1. **Envoy:** *proxy*, to *mediate* all inbound and outbound traffic for all services in the service mesh
2. **Pilot:** *Push configuration to envoy fleet*, responsible for service discovery, registration and load balancing
3. **Istio-Auth** provides strong *service-to-service* and end-user authentication using *mutual TLS*
4. **Mixer** is responsible for enforcing *access control* across the service mesh & collecting *telemetry* data from the Envoy proxy and other services




How it works

Life of a request in the mesh

Service A comes up.
Envoy is deployed
alongside it and fetches
service information,
routing and configuration
policy from Pilot.




Life of a request in the mesh




Service A makes a call to service B

Client-side Envoy intercepts the call.

Envoy consults config from Pilot to know how/where to route call to service B




Service A




Service B

Life of a request in the mesh




Envoy forwards request
to appropriate instance of
service B




Life of a request in the mesh

Server-side Envoy checks with Mixer to validate that call should be allowed (ACL check, quota check, etc).




Life of a request in the mesh




Mixer checks with appropriate adaptors (policy engine, quota adaptor) to verify that the call can proceed


Life of a request in the mesh



Server-side Envoy forwards request to service B, which processes the request and returns response.




Service A




Service B


Life of a request in the mesh





Envoy forwards response
to the caller




Life of a request in the mesh



Client-side Envoy forwards response to original caller.




Life of a request in the mesh



Envoy reports telemetry to Mixer, which in turn notifies appropriate plugins

Life of a request in the mesh



Client-side Envoy reports telemetry to Mixer (including client-perceived latency), which in turn notifies appropriate plugins

How to “operate” Istio

- Dedicated command line tool
 - [Istioctl](#)
 - `istioctl kube-inject`

Example configuration

Deploy BookInfo sample on Kubernetes:

```
kubectl apply -f <(istioctl kube-inject -f samples/bookinfo/kube/bookinfo.yaml)
```


Example configuration

Create a custom RouteRule

```
cat <<EOF | kubectl create -f -  
  
apiVersion: config.istio.io/v1alpha2  
kind: RouteRule  
metadata:  
  name: reviews-test-v2  
  namespace: default  
  ...  
spec:  
  destination:  
    name: reviews  
  match:  
    request:  
      headers:  
        cookie:  
          regex: ^(*?;)?(user=jason)(;.*)?$  
  precedence: 2  
  route:  
  - labels:  
    version: v2
```

EOF

Easiest way to play with Istio

<https://www.katacoda.com/rafabene/courses/workshop/>

Thanks!

Twitter: @geoand86

Email: geoand@gmail.com

Slides: <http://slideshare.net/GeorgiosAndrianakis>

Further reading

- <http://envoyproxy.io>
- <http://istio.io>
- <https://medium.com/@mattklein123/>
- <http://blog.christianposta.com/istio-workshop/slides/>
- <http://blog.christianposta.com/tags/#istio>

Q & A

