



Transactional Outbox

with a twist

Chris Aslanoglou - Software Engineer
Java Hellenic User Group (JHUG) - December 2022

The problem

An example

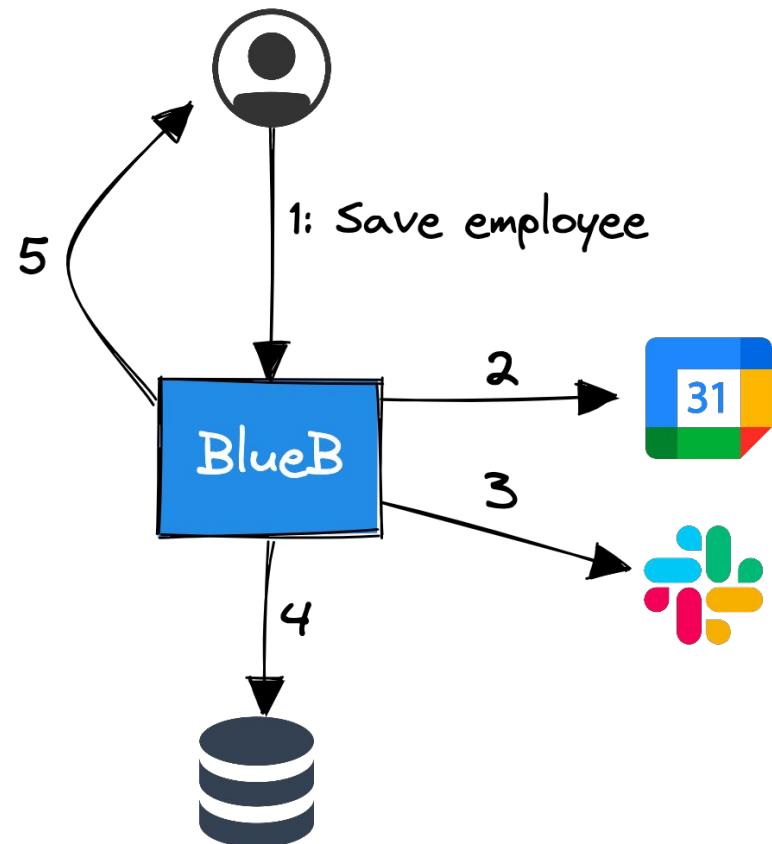
Goal

1. Save an employee's details
2. Create new GCalendar
3. Fetch SlackId

What can go wrong? famous last words

- Transient network failures Google
- Same with Slack
- DB consistency failure => rollback
- It's slow

Employee save flow



Outbox pattern to the rescue

Outbox Pattern

Q: What do we want to achieve?

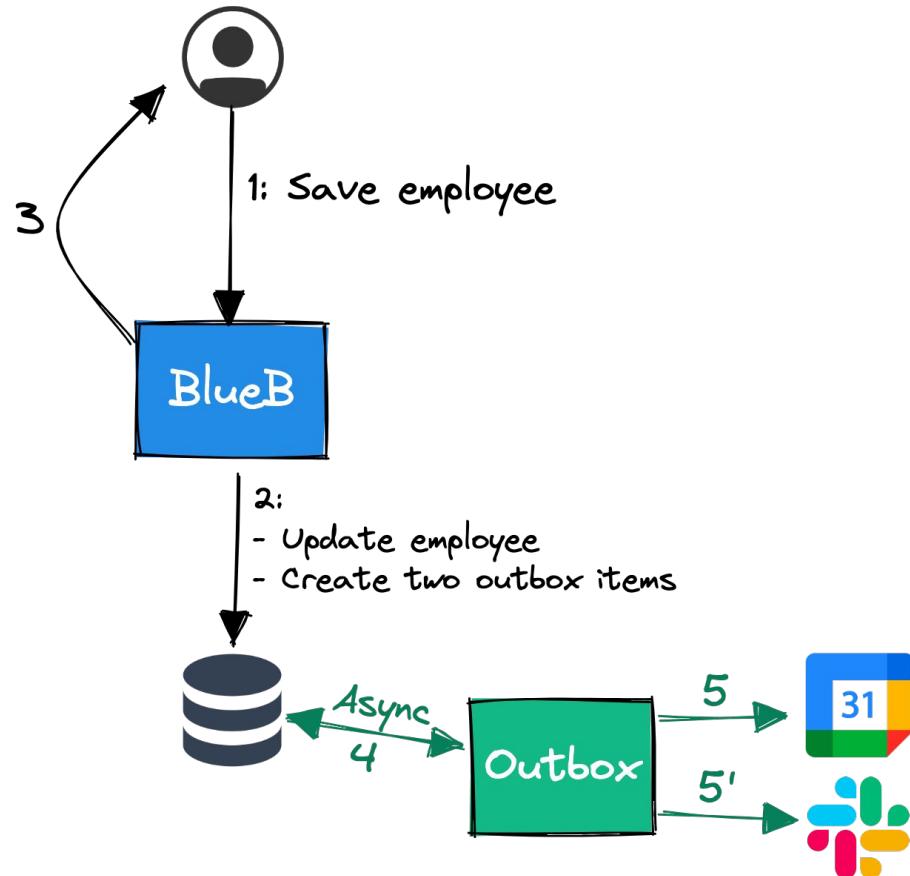
A: Have our cake and eventually eat it

Well, kinda...

Reliably and atomically
update a DB record &
communicate with another service

Create a new Google calendar
if and only if
the “save employee” transaction succeeds

Employee save flow



Our requirements

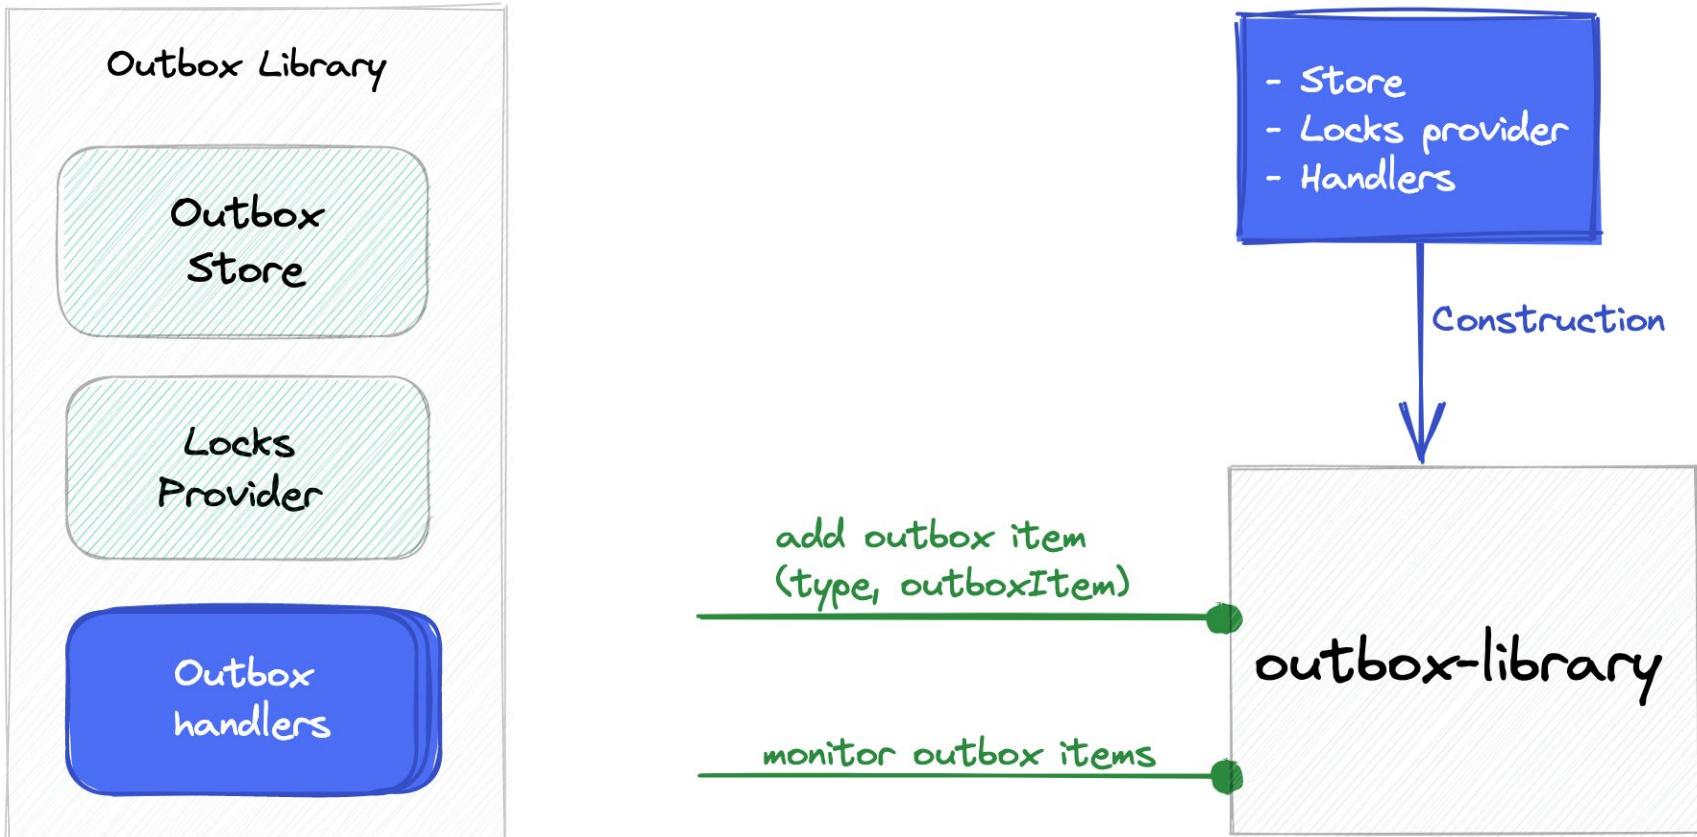
Our requirements

- A library to persist and handle generic outbox items
- At-least once processing semantics
- Per type failure handling granularity
- Agnostic to
 - Application framework
 - Data storage layer
 - Locks provider implementation

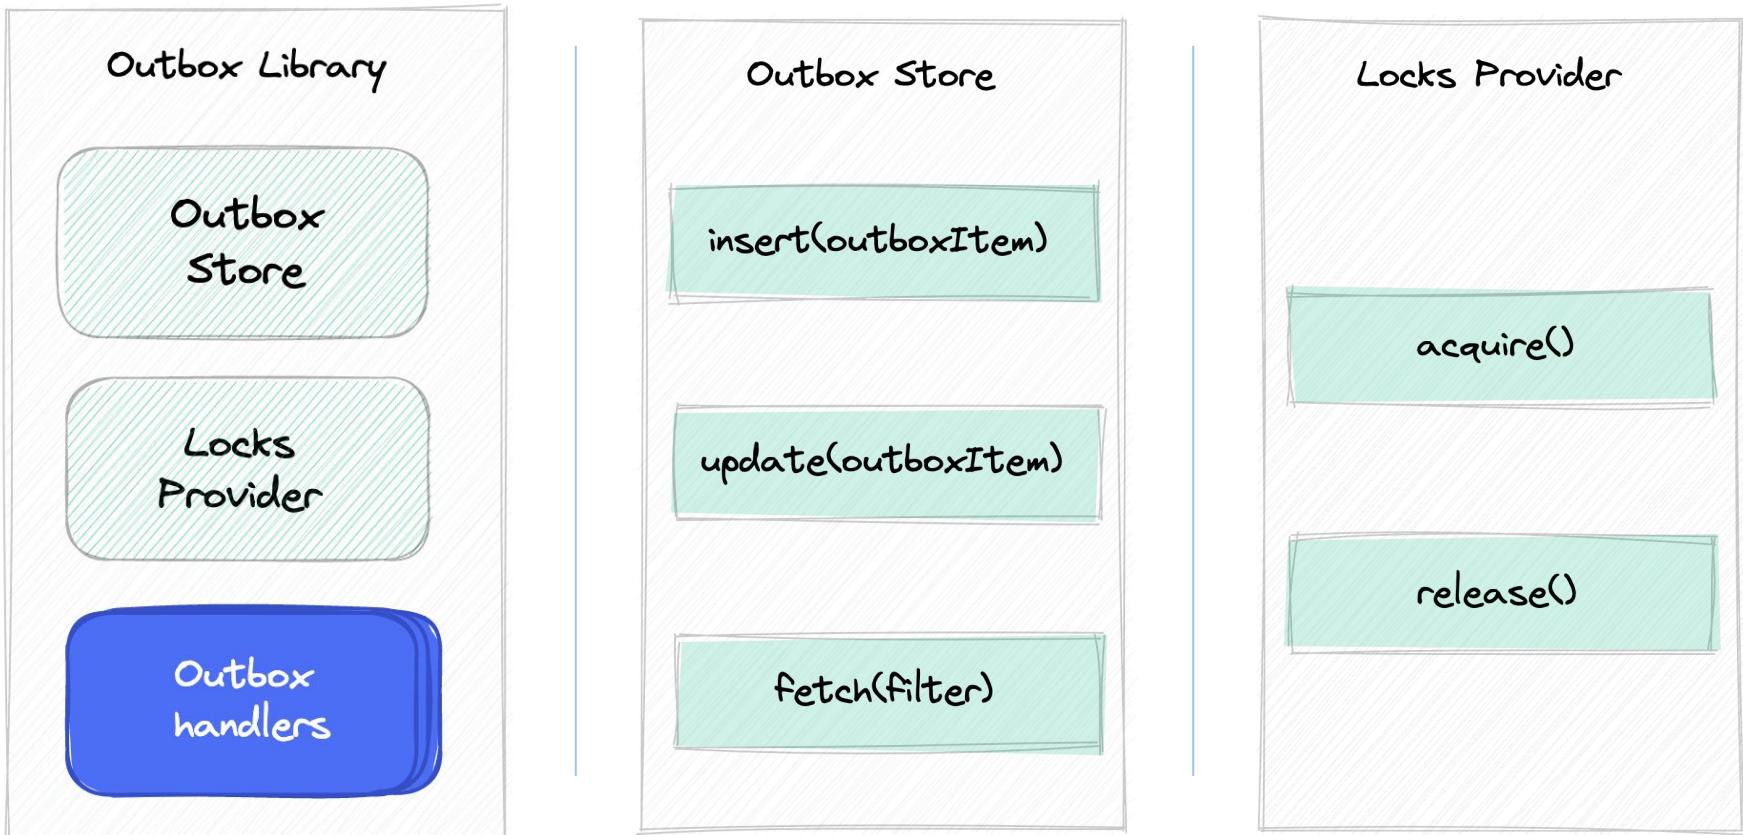


Our approach

High Level Design - Core Library



High Level Design - Required Interfaces



Handlers, Types, Items, or, how do I use this thing?

Outbox Item

- id
- type
- serial. payload
- status
 - pending
 - running
 - completed
 - failed
- retries
- nextRun

Outbox Handler

- getSupportedType()
- serialize(payload)
- handle(serializedPayload)
- hasReachedMaxRetries(retries)
- getNextExecutionTime(currentRetries)
- handleFailure(serializedPayload)

Library instantiation

```
return TransactionalOutboxBuilder  
    .make(clock)  
    .withHandlers(outboxHandlers)  
    .withLocksProvider(locksProvider)  
    .withOutboxStore(outboxStore)  
    .build()
```

Item addition

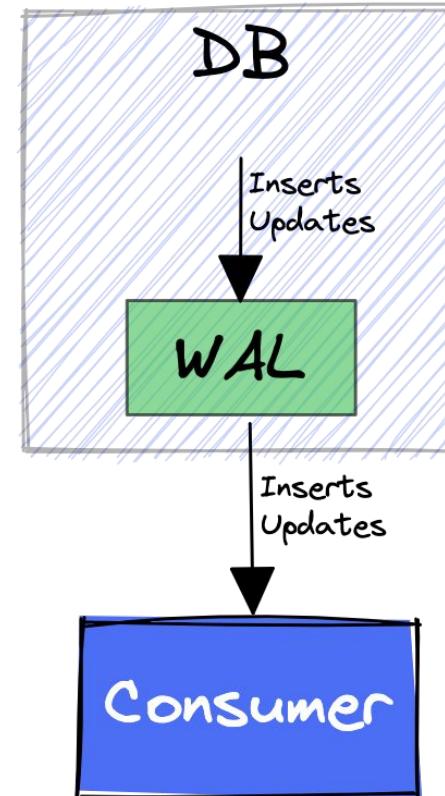
```
private void addGoogleCalendarOutboxItem(User user) {  
    var payload = new GoogleCalendarCreationOutboxPayload(  
        user.getId(),  
        user.getFullname()  
    );  
    outbox.add(GOOGLE_CALENDAR_CREATE, payload);  
}
```

Ok, we saved the items, now what?

Polling Consumer



Change-Data-Capture



Polling consumer call

```
@Scheduled(cron = "\${app.outbox.cron}")
@Transactional
fun run() {
    outbox.monitor()
}
```

Polling Consumer Outline

```
1. fun monitor() {  
2.     locksProvider.acquire()  
3.  
4.     val items = store.fetch(eligibleFilter)  
5.  
6.     markForProcessing(items)  
7.     items.map { store.update(it) }  
8.  
9.     items.forEach { item ->  
10.         executor.execute(  
11.             OutboxItemProcessor(item, handlers[item.type], store)  
12.         )  
13.     }  
14.  
15.     locksProvider.release()  
16. }
```

Polling Consumer Outline

```
1. fun monitor() {  
2.     locksProvider.acquire()  
3.  
4.     val items = store.fetch(eligibleFilter)  
5.  
6.     markForProcessing(items)  
7.     items.map { store.update(it) }  
8.  
9.     items.forEach {  
10.         executor.execute(  
11.             OutboxItemPr  
12.         )  
13.     }  
14.  
15.     locksProvider.re  
16. }
```

```
1. fun markForProcessing(items) {  
2.     items.map {  
3.         it.status = OutboxStatus.RUNNING  
4.         val now = Instant.now(clock)  
5.         it.lastExecution = now  
6.         it.rerunAfter = now.plus(rerunAfterDuration)  
7.     }  
8. }
```

OutboxItemProcessor (thread) Outline

```
1. fun run() {
2.     try {
3.         handler.handle(item.payload)
4.         item.status = OutboxStatus.COMPLETED
5.
6.     } catch (exception: Exception) {
7.         if (handler.hasReachedMaxRetries(item.retries)) {
8.             handleTerminalFailure(exception)
9.         }
10.    else {
11.        handleRetryableFailure(exception)
12.    }
13. } finally {
14.     store.update(item)
15. }
16. }

// handleTerminalFailure => item.status = FAILED && handler.handleFailure(item.payload)

// handleRetryableFailure =>
//   1. item.nextRun = handler.getNextExecutionx(item.retries)
//   2. item.retries += 1
//   3. item.status = OutboxStatus.PENDING
```

Required API implementation

Store

- Simple data access layer abstraction
- **Insert**
TX semantics of outbox.add caller
- **Fetch**
Read only TX
- **Update**
Requires new TX
 - Outbox item update committed before thread execution
=> Avoid concurrent updates
 - In thread: always commit in new TX

Locks Provider

- Requirements
 - Blocking exclusive locks
 - Automatic release (crash-failure)
- Implemented with PostgreSQL advisory locks
 - Lock semantics prescribed by application
 - DB session connection scoped
 - Released after connection is closed (crash-failure liveness)
 - Blocking

Acquire: `select pg_advisory_lock(:id)`

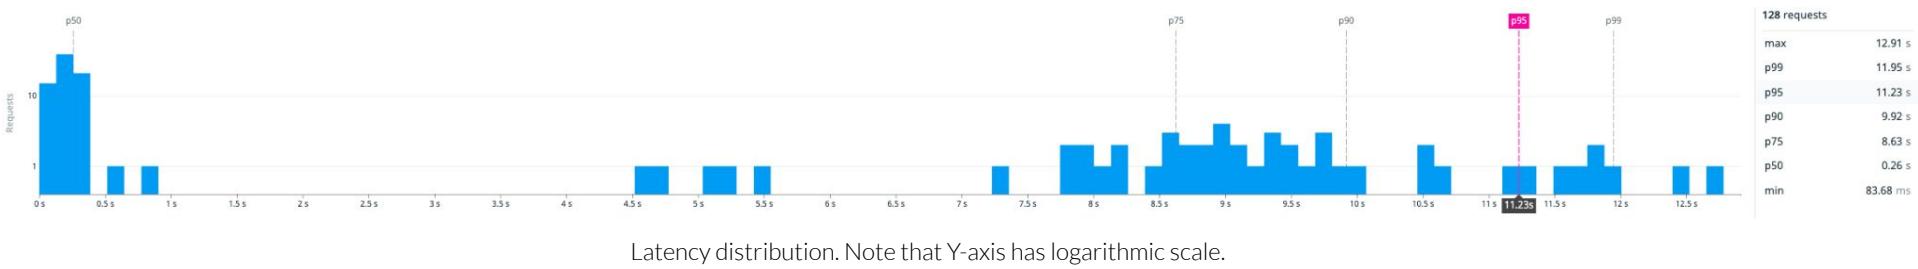
Release: `select pg_advisory_unlock(:id)`



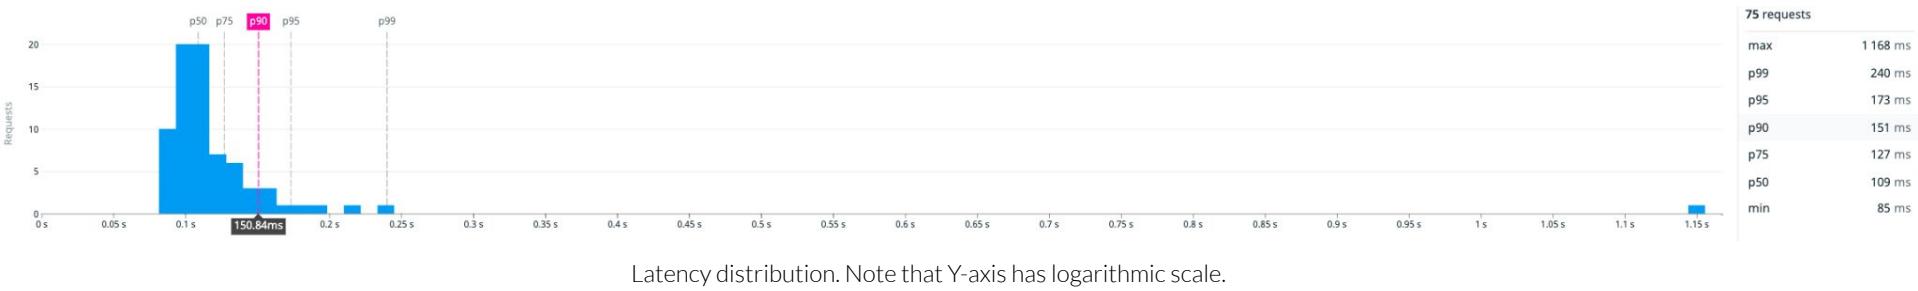
GAINZ

User update latencies

Historic data before outbox release



Historic data after the release



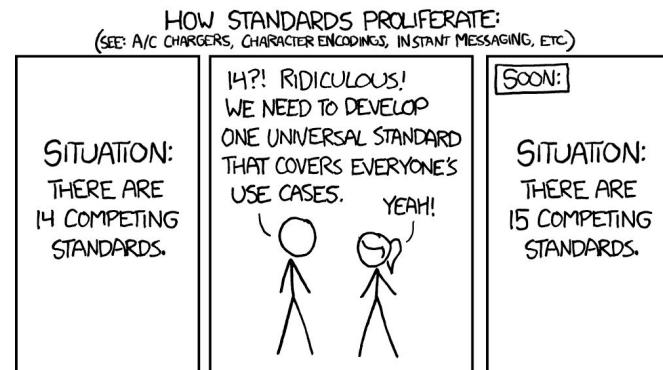
What we've gained

- Reduced boilerplate code for new outbox items handling
 - Handler delegates to existing business flows
- Retries out of the box
- Per-type failure handling
- Learned much along the way



Related work

- Eventuate Tram (Richardson, microservices.io), quite thorough solution
 - More moving parts
 - message broker
 - their Change-Data-Capture microservice: [requires Zookeeper for leader election when running as a cluster for high availability](#)
- gruelbox/transaction-outbox, great alternative
 - Missing flexibility on a per-outbox type handlers (i.e., for failure, retry strategies) (e.g., they provide a single failure handler)
- Also, XKCD-927 and above all a learning opportunity



Concluding

Future work

- Configurable retry strategies (linear, exponential back-off)
- Immediate outbox item handling (e.g., via a post-commit TX hook or a CDC consumer)
- Evaluate co-routines instead of Executor service
- Extract to a library
 - core
 - spring-outbox: wrapper on-top core with store and locks implementations

Team

- Thanasis Polydoros
- Apostolis Kiraleos

Thank you